

0.) Import and Clean data

```
In [1]: import pandas as pd
        from google.colab import drive
        import matplotlib.pyplot as plt
        import numpy as np
```

```
In [2]: from sklearn.linear_model import LogisticRegression
        from sklearn.tree import DecisionTreeClassifier
        from sklearn.ensemble import BaggingClassifier
        from sklearn.datasets import make_classification
        from sklearn.metrics import accuracy_score
        from sklearn.model_selection import train_test_split
        from sklearn.preprocessing import StandardScaler
        from sklearn.tree import plot_tree
        from sklearn.metrics import confusion_matrix
        import seaborn as sns
```

```
In [ ]: #drive.mount('/content/gdrive/', force_remount = True)
```

Mounted at /content/gdrive/

```
In [4]: df = pd.read_csv("bank-additional-full.csv", sep = ";")
```

```
In [5]: df.head()
```

```
Out[5]:
```

	age	job	marital	education	default	housing	loan	contact	month	day_of_week	...	campaign	pday
0	56	housemaid	married	basic.4y	no	no	no	telephone	may	mon	...	1	99
1	57	services	married	high.school	unknown		no	telephone	may	mon	...	1	99
2	37	services	married	high.school	no	yes	no	telephone	may	mon	...	1	99
3	40	admin.	married	basic.6y	no	no	no	telephone	may	mon	...	1	99
4	56	services	married	high.school	no	no	yes	telephone	may	mon	...	1	99

5 rows × 21 columns

```
In [6]: df = df.drop(["default", "pdays", "previous", "poutcome", "emp.var.rate",
                    "cons.price.idx", "cons.conf.idx", "euribor3m", "nr.employed"], axis =
                    1)
        df = pd.get_dummies(df, columns = ["loan",
        "job", "marital", "housing", "contact", "day_of_week", "campaign", "month",
        "education"], drop_first = True)
```

```
In [7]: df.head()
```


```
Out[7]:
```

	age	duration	y	loan_unknown	loan_yes	job_blue-collar	job_entrepreneur	job_housemaid	job_management	job_reti
0	56	261	no	0	0	0	0	1	0	
1	57	149	no	0	0	0	0	0	0	
2	37	226	no	0	0	0	0	0	0	
3	40	151	no	0	0	0	0	0	0	
4	56	307	no	0	1	0	0	0	0	

5 rows × 83 columns

```
In [8]: y = pd.get_dummies(df["y"], drop_first = True)
X = df.drop(["y"], axis = 1)
```

```
In [9]: obs = len(y)
plt.bar(["No", "Yes"], [len(y[y. yes==0])/obs, len(y[y. yes==1])/obs])
plt.ylabel("Percentage of Data")
plt.show()
```



```
In [10]: # Train Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
random_state=42)

scaler = StandardScaler().fit(X_train)

X_scaled = scaler.transform(X_train)
X_test = scaler.transform(X_test)
```

1.) Based on the visualization above, use your expert opinion to transform the data based on what we

learned this quarter

```
In [54]: #####  
#####TRANSFORM###  
#####  
  
# Transform by SMOTE  
from imblearn.over_sampling import SMOTE  
  
smote = SMOTE()  
smote_X, smote_y = smote.fit_resample(X_train, y_train)
```

```
In [55]: X_scaled = smote_X  
y_train = smote_y  
X_scaled  
y_train
```

```
Out[55]:
```

	yes
0	0
1	0
2	0
3	0
4	0
...	...
51155	1
51156	1
51157	1
51158	1
51159	1

51160 rows × 1 columns

```
In [58]: len(X_scaled)
```

```
Out[58]: 51160
```

```
In [59]: len(y_train)
```

```
Out[59]: 51160
```

2.) Build and visualize a decision tree of Max Depth 3. Show the confusion matrix.


```
In [26]: dtree_main = DecisionTreeClassifier(max_depth = 3)  
dtree_main.fit(X_scaled, y_train)
```

Out[26]: ▾ DecisionTreeClassifier
DecisionTreeClassifier(max_depth=3)

```
In [27]: fig, axes = plt.subplots(nrows = 1,ncols = 1,figsize = (4,4), dpi=300)
plot_tree(dtree_main, filled = True, feature_names = X.columns, class_names=
["No", "Yes"])

#fig.savefig('imagename.png')
```

```
Out[27]: [Text(0.5, 0.875, 'duration <= 1.025\ngini = 0.2\nsamples = 28831\nvalue = [25580, 3251]\nclass = N
o'),
Text(0.25, 0.625, 'duration <= -0.197\ngini = 0.134\nsamples = 25702\nvalue = [23841, 1861]\nclass =
No'),
Text(0.125, 0.375, 'month_mar <= 4.161\ngini = 0.063\nsamples = 16326\nvalue = [15795, 531]\nclass =
No'),
Text(0.0625, 0.125, 'gini = 0.055\nsamples = 16102\nvalue = [15644, 458]\nclass = No'),
Text(0.1875, 0.125, 'gini = 0.439\nsamples = 224\nvalue = [151, 73]\nclass = No'),
Text(0.375, 0.375, 'month_oct <= 3.6\ngini = 0.243\nsamples = 9376\nvalue = [8046, 1330]\nclass = N
o'),
Text(0.3125, 0.125, 'gini = 0.227\nsamples = 9175\nvalue = [7980, 1195]\nclass = No'),
Text(0.4375, 0.125, 'gini = 0.441\nsamples = 201\nvalue = [66, 135]\nclass = Yes'),
Text(0.75, 0.625, 'duration <= 2.22\ngini = 0.494\nsamples = 3129\nvalue = [1739, 1390]\nclass = N
o'),
Text(0.625, 0.375, 'contact_telephone <= 0.283\ngini = 0.466\nsamples = 2028\nvalue = [1279, 749]\ncl
ass = No'),
Text(0.5625, 0.125, 'gini = 0.49\nsamples = 1364\nvalue = [777, 587]\nclass = No'),
Text(0.6875, 0.125, 'gini = 0.369\nsamples = 664\nvalue = [502, 162]\nclass = No'),
Text(0.875, 0.375, 'contact_telephone <= 0.283\ngini = 0.486\nsamples = 1101\nvalue = [460, 641]\ncla
ss = Yes'),
Text(0.8125, 0.125, 'gini = 0.47\nsamples = 745\nvalue = [281, 464]\nclass = Yes'),
Text(0.9375, 0.125, 'gini = 0.5\nsamples = 356\nvalue = [179, 177]\nclass = No')]
```




1b.) Confusion matrix on out of sample data. Visualize and store as variable

```
In [29]: y_pred = dtree_main.predict(X_test)
y_true = y_test
cm_raw = confusion_matrix(y_true, y_pred)
```

```
In [30]: class_labels = ['Negative', 'Positive']

# Plot the confusion matrix as a heatmap
sns.heatmap(cm_raw, annot=True, fmt='d', cmap='Blues', xticklabels=class_labels,
yticklabels=class_labels)
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```



3.) Use bagging on your descision tree

```
In [31]: #placeholder for optimizing max depth
dtree = DecisionTreeClassifier(max_depth = 3)
```

```
In [32]: bagging = BaggingClassifier(estimator = dtree,
                                     n_estimators = 100,
                                     max_samples = .5,
                                     max_features = 1.)
```

```
In [33]: bagging.fit(X_scaled, y_train)

y_pred = bagging.predict(X_test)
```


```
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_bagging.py:802: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
```

```
In [34]: y_true = y_test
cm_raw = confusion_matrix(y_true, y_pred)
```

```
In [35]: class_labels = ['Negative', 'Positive']

# Plot the confusion matrix as a heatmap
sns.heatmap(cm_raw, annot=True, fmt='d', cmap='Blues', xticklabels=class_labels,
            yticklabels=class_labels)
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
```

```
plt.ylabel('True Label')
plt.show()
```



4.) Boost your tree

```
In [36]: from sklearn.ensemble import AdaBoostClassifier
```

```
In [37]: #placeholder for optimizing max depth
dtree = DecisionTreeClassifier(max_depth = 3)
```

```
In [38]: boost = AdaBoostClassifier(estimator = dtree,
                                   n_estimators = 100,
                                   learning_rate = .1)
```

```
In [39]: boost.fit(X_scaled, y_train)


y_pred = boost.predict(X_test)

y_true = y_test
cm_raw = confusion_matrix(y_true, y_pred)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
```

```
In [40]: # Plot the confusion matrix as a heatmap
sns.heatmap(cm_raw, annot=True, fmt='d', cmap='Blues', xticklabels=class_labels,
            yticklabels=class_labels)
```

```
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```



5.) Train a Logistic Regression (Super Learner) on the Decision tree, Boosted tree, Bagged tree.

Interpret coefficients and significance

```
In [ ]: # pip install mlens
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: mlens in /usr/local/lib/python3.8/dist-packages (0.2.3)
Requirement already satisfied: scipy>=0.17 in /usr/local/lib/python3.8/dist-packages (from mlens) (1.10.1)
Requirement already satisfied: numpy>=1.11 in /usr/local/lib/python3.8/dist-packages (from mlens) (1.22.4)
```

```
In [41]: from sklearn.linear_model import LogisticRegression
```

```
In [42]: base_predictions = [dtree_main.predict(X_train),
                             boost.predict(X_train),
                             bagging.predict(X_train)]
```



```

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432: UserWarning: X has feature names, but DecisionTreeClassifier was fitted without feature names
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432: UserWarning: X has feature names, but AdaBoostClassifier was fitted without feature names
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432: UserWarning: X has feature names, but BaggingClassifier was fitted without feature names
  warnings.warn(

```

In [43]: `base_predictions`

Out[43]: `[array([1, 0, 0, ..., 1, 0, 1], dtype=uint8),
array([1, 1, 1, ..., 1, 1, 1], dtype=uint8),
array([1, 1, 1, ..., 1, 1, 1], dtype=uint8)]`

In [44]: `base_predictions = [list(dtree_main.predict(X_train)),
list(boost.predict(X_train)),
list(bagging.predict(X_train))]`

```

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432: UserWarning: X has feature names, but DecisionTreeClassifier was fitted without feature names
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432: UserWarning: X has feature names, but AdaBoostClassifier was fitted without feature names
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432: UserWarning: X has feature names, but BaggingClassifier was fitted without feature names
  warnings.warn(

```

In [47]: `np.array(base_predictions)`

Out[47]: `array([[1, 0, 0, ..., 1, 0, 1],
[1, 1, 1, ..., 1, 1, 1],
[1, 1, 1, ..., 1, 1, 1]], dtype=uint8)`

In [48]: `n = len(base_predictions[0])`

In [50]: `base_predictions = [np.array(base_predictions)[:i] for i in range(n)]`

In [51]: `super_learner = LogisticRegression()`

In [52]: `super_learner.fit(base_predictions, y_train)`

```

/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)

```

Out[52]: `▼ LogisticRegression
LogisticRegression()`

In [53]: `super_learner.coef_`

Out[53]: `array([[1.17918485, 0.11003523, 0.23200188]])`

Based on the results, it seems like the Decision Tree is exactly the best.

And it weights $1.17 \times \hat{y}_{DecisionTree} + 0.11 \times \hat{y}_{BoostedTree} + 0.23 \times \hat{y}_{BaggedTree}$.

