

0.) Import the Credit Card Fraud Data From CCLE

```
In [2]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import warnings
warnings.filterwarnings("ignore")
```

```
In [3]: #drive.mount('/content/gdrive/', force_remount = True)
```

```
In [4]: df = pd.read_csv("fraudTest.csv")
```

```
In [5]: df.head()
```

```
Out[5]:
```

	Unnamed: 0	trans_date_trans_time	cc_num	merchant	category	amt	first	last	gender	street	...	lat	long	city_pop
0	0	2020-06-21 12:14:25	2291163933867244	fraud_Kirlin and Sons	personal_care	2.86	Jeff	Elliott	M	351 Darlene Green	...	33.9659	-80.9355	3334
1	1	2020-06-21 12:14:33	3573030041201292	fraud_Sporer-Keebler	personal_care	29.84	Joanne	Williams	F	3638 Marsh Union	...	40.3207	-110.4360	3
2	2	2020-06-21 12:14:53	3598215285024754	fraud_Swaniawski, Nitzsche and Welch	health_fitness	41.28	Ashley	Lopez	F	9333 Valentine Point	...	40.6729	-73.5365	34
3	3	2020-06-21 12:15:15	3591919803438423	fraud_Haley Group	misc_pos	60.05	Brian	Williams	M	32941 Krystal Mill Apt. 552	...	28.5697	-80.8191	54
4	4	2020-06-21 12:15:17	3526826139003047	fraud_Johnston-Casper	travel	3.19	Nathan	Massey	M	5783 Evan Roads Apt. 465	...	44.2529	-85.0170	1

5 rows × 23 columns

```
In [6]: df_select = df[["trans_date_trans_time", "category", "amt", "city_pop", "is_fraud"]]

df_select["trans_date_trans_time"] = pd.to_datetime(df_select["trans_date_trans_time"])
df_select["time_var"] = [i.second for i in df_select["trans_date_trans_time"]]

X = pd.get_dummies(df_select, ["category"]).drop(["trans_date_trans_time", "is_fraud"], axis = 1)
y = df["is_fraud"]
```

1.) Use scikit learn preprocessing to split the data into 70/30 in out of sample

```
In [7]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
```

```
In [8]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = .3)
```

```
In [9]: X_test, X_holdout, y_test, y_holdout = train_test_split(X_test, y_test, test_size = .5)
```

```
In [10]: X_train_df = X_train
X_test_df = X_test
X_holdout_df = X_holdout
```

```
In [11]: scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
X_holdout = scaler.transform(X_holdout)
```

2.) Make three sets of training data (Oversample, Undersample and SMOTE)

```
In [12]: from imblearn.over_sampling import RandomOverSampler
from imblearn.under_sampling import RandomUnderSampler
from imblearn.over_sampling import SMOTE
```

```
In [13]: ros = RandomOverSampler()
over_X, over_y = ros.fit_resample(X_train, y_train)

rus = RandomUnderSampler()
under_X, under_y = rus.fit_resample(X_train, y_train)

smote = SMOTE()
smote_X, smote_y = smote.fit_resample(X_train, y_train)
```

```
In [14]: len(y_train)
```

```
Out[14]: 389003
```

```
In [15]: sum(y_train ==1)
```

```
Out[15]: 1530
```

```
In [16]: sum(y_train ==0)
```

```
Out[16]: 387473
```

```
In [17]: len(over_y)
```

```
Out[17]: 774946
```

```
In [18]: len(under_y)
```

```
Out[18]: 3060
```

```
In [19]: len(smote_y)
```

```
Out[19]: 774946
```

3.) Train three logistic regression models

```
In [20]: from sklearn.linear_model import LogisticRegression
```

```
In [21]: over_log = LogisticRegression().fit(over_X, over_y)

under_log = LogisticRegression().fit(under_X, under_y)

smote_log = LogisticRegression().fit(smote_X, smote_y)
```

4.) Test the three models

```
In [22]: over_log.score(X_test, y_test)
```

```
Out[22]: 0.91798027783776
```

```
In [23]: under_log.score(X_test, y_test)
```

```
Out[23]: 0.9225029391300175
```

```
In [24]: smote_log.score(X_test, y_test)
```

```
Out[24]: 0.9176323808152786
```

```
In [25]: # We see SMOTE performing with higher accuracy but is ACCURACY really the best measure?
```

5.) Which performed best in Out of Sample metrics?

```
In [26]: # Sensitivity here in credit fraud is more important as seen from last class
```

```
In [27]: from sklearn.metrics import confusion_matrix
```

```
In [28]: y_true = y_test
```

```
In [29]: y_pred = over_log.predict(X_test)
cm = confusion_matrix(y_true, y_pred)
cm
```

```
Out[29]: array([[76292,  6750],
               [   87,   229]], dtype=int64)
```

```
In [30]: print("Over Sample Sensitivity : ", cm[1,1] / ( cm[1,0] + cm[1,1]))
```

Over Sample Sensitivity : 0.7246835443037974

```
In [31]: y_pred = under_log.predict(X_test)
cm = confusion_matrix(y_true, y_pred)
cm
```

```
Out[31]: array([[76670,  6372],
               [   88,   228]], dtype=int64)
```

```
In [32]: print("Under Sample Sensitivity : ", cm[1,1] / ( cm[1,0] + cm[1,1]))
```

Under Sample Sensitivity : 0.7215189873417721

```
In [33]: y_pred = smote_log.predict(X_test)
cm = confusion_matrix(y_true, y_pred)
cm
```

```
Out[33]: array([[76263,  6779],
               [   87,   229]], dtype=int64)
```

```
In [34]: print("SMOTE Sample Sensitivity : ", cm[1,1] / ( cm[1,0] + cm[1,1]))
```

SMOTE Sample Sensitivity : 0.7246835443037974


6.) Pick two features and plot the two classes before and after SMOTE.

```
In [35]: raw_temp = pd.concat([X_train_df, y_train], axis =1)
```

```
In [36]: #plt.scatter(raw_temp[raw_temp["is_fraud"] == 0]["amt"], raw_temp[raw_temp["is_fraud"] == 0]["city_pop"])
```

```
plt.scatter(raw_temp[raw_temp["is_fraud"] == 1]["amt"], raw_temp[raw_temp["is_fraud"] == 1]["city_pop"])
plt.legend(["Fraud", "Not Fraud"])
plt.xlabel("Amount")
plt.ylabel("Population")
```

```
plt.show()
```




```
In [37]: smote_X_df, smote_y = smote.fit_resample(X_train_df, y_train)
raw_temp = pd.concat([smote_X_df, smote_y], axis =1)
```

```
In [38]: #plt.scatter(raw_temp[raw_temp["is_fraud"] == 0]["amt"], raw_temp[raw_temp["is_fraud"] == 0]["city_pop"])

plt.scatter(raw_temp[raw_temp["is_fraud"] == 1]["amt"], raw_temp[raw_temp["is_fraud"] == 1]["city_pop"])
plt.legend(["Not Fraud", "Fraud"])
plt.xlabel("Amount")
plt.ylabel("Population")

plt.show()
```



7.) We want to compare oversampling, Undersampling and SMOTE across our 3 models (Logistic Regression, Logistic Regression Lasso and Decision Trees).

Make a dataframe that has a dual index and 9 Rows.

Calculate: Sensitivity, Specificity, Precision, Recall and F1 score. for out of sample data.

Notice any patterns across performance for this model. Does one totally outperform the others IE. over/under/smote or does a model perform better DT, Lasso, LR?

Choose what you think is the best model and why. test on Holdout

```
In [39]: from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score
import pandas as pd
```

```
In [40]: resampling_methods = {
    "over": RandomOverSampler(),
    "under": RandomUnderSampler(),
    "smote": SMOTE()
}

model_configs = {
    "LOG": LogisticRegression(),
    "LASSO": LogisticRegression(penalty = "l1", solver = "liblinear", C = .5),
    "DecisionTree": DecisionTreeClassifier()
}
```

```
In [41]: trained_models = {}
results = []
```

```
In [42]: def calc_perf_metric(y_true, y_pred):
        tn, fp, fn, tp = confusion_matrix(y_true, y_pred).ravel()

        sensitivity = tp/(tp+fn)
        specificity = tn/(tn+fp)
        precision = precision_score(y_true, y_pred)
        recall = recall_score(y_true, y_pred)
        f1 = f1_score(y_true, y_pred)

        return(sensitivity, specificity, precision, recall, f1)
```

```
In [43]: for resample_key, resampler in resampling_methods.items():
        resample_X, resample_y = resampler.fit_resample(X_train, y_train)

        for model_name, model in model_configs.items():

            combined_key = f"{resample_key}_{model_name}"
            m = model.fit(resample_X, resample_y)
            trained_models[combined_key] = m
            y_pred=m.predict(X_test)
            sensitivity, specificity, precision, recall, f1=calc_perf_metric(y_test, y_pred)
            results.append({"Model":combined_key,
                           "Sensitivity":sensitivity,
                           "Specificity":specificity,
                           "Precision":precision,
                           "Recall":recall,
                           "F1":f1})
```

```
In [44]: results_df = pd.DataFrame(results)
        results_df
```

Out[44]:

	Model	Sensitivity	Specificity	Precision	Recall	F1
0	over_LOG	0.724684	0.917825	0.032468	0.724684	0.062152
1	over_LASSO	0.724684	0.917849	0.032478	0.724684	0.062169
2	over_DecisionTree	0.563291	0.998531	0.593333	0.563291	0.577922
3	under_LOG	0.702532	0.935960	0.040072	0.702532	0.075820
4	under_LASSO	0.702532	0.935996	0.040094	0.702532	0.075859
5	under_DecisionTree	0.955696	0.943547	0.060521	0.955696	0.113833
6	smote_LOG	0.724684	0.917283	0.032263	0.724684	0.061775
7	smote_LASSO	0.724684	0.917295	0.032267	0.724684	0.061783
8	smote_DecisionTree	0.718354	0.993618	0.299868	0.718354	0.423113